


Welcome to Lecture 38!

Topic: Databases

Databases: The Mental Model

What a database is, what it solves, and how to choose SQL vs NoSQL



Question: If your server restarts right now, what happens to user data?

Why do we need a database?

Attempt 1: In memory

Data in variables / arrays.



Fast



Restart wipes everything



“Deploy → data = gone”

Attempt 2: A file (JSON/CSV)

- Survives restart
- Searching = scan whole file
- Two writes can clash
- Partial write can corrupt data



Attempt 3: Database

- Persistent
- Fast search
- Safe concurrent updates
- Recovery tools




Databases exist because apps need data to be: persistent, searchable, and correct.

What is a Database?

Database vs DBMS (the software that runs it)

Database



An organized store of data built for reliable reads/writes and fast questions (queries).

DBMS (Database Management System)

The software that stores, indexes, secures, and serves your data.

Examples: MySQL, PostgreSQL, MongoDB.

Think: DB = data • DBMS = engine

What does a Database give you?

● Persistence

Data survives restarts/crashes.

● Querying

Ask questions like: “marks > 70”.

● Indexing

Fast lookup without scanning everything.

● Concurrency

Many users can write safely.

● Transactions

All-or-nothing updates.

● Constraints

Keep bad data out (unique, not null).

● Backup & recovery

Undo accidents / restore after failure.

Tip: when choosing a *type* of database for different applications, you're mostly choosing trade-offs across these features.

Persistence

"If the server restarts... do we still have the data?"

In-memory

Fast
But disappears on restart



File

Survives restart
But hard to update safely



Database

Survives restart
+ safe updates



Think like an engineer:

"If a request is in-flight and the server crashes mid-write, what state is the data in?"

Querying

“Ask questions about data without writing loops every time”

Typical app questions:

- “Find student whose id=42”
- “Students with marks > 70”
- “Orders from last 7 days”
- “Top 5 best-selling products”

Key idea:

Your app asks a question and DB returns matching rows/documents.

Without a DB (file scan):

for each record:
if matches → keep it

With a DB:

Execute query -> results



Indexing

A fast lookup structure that avoids scanning everything

Analogy: book index

Without an index:

You flip every page until you find the word.

With an index:

You jump directly to the page.

Ex: Finding names in a phonebook

Asha -> p. 12

Sucheta -> p. 41

Zoya -> p. 58

Conceptual performance

Scan the whole file



Index lookup



Takeaway:

Indexes speed up reads, but add work on writes (the index must be updated).

Concurrency Control

When multiple users read/write at the same time

Scenario: “Last seat available”

Two users click “Book” at the same time.

Problem: both users book 1A at the same time will lead to double booking

User A

View seats

Decide to book seat 1A

Book seat 1A

Confirm booking

User B

View seats

Decide to book seat 1A

Book seat 1A

Confirm booking

Databases provide safe concurrent updates so you don’t “lose” writes.

Transactions

All-or-nothing updates (to prevent half-finished state)

Scenario: placing an order

These steps must stay consistent:

- 1) Create order record
- 2) Reduce inventory
- 3) Record payment

What can go wrong?

Order created ✓
Inventory reduced ✓
Payment recorded ✗ (crash)

Transaction guarantee:

Either all 3 happen, or none happen.

In one sentence: A transaction prevents the database from ending up in a “half-updated” state.

Different ways to store data

Same information, different shapes

Tabular (tables)

Students		
id	name	mark
1	Asha	85
2	Rahul	62

Great for structured data + relationships.

Eg: MySQL, Postgres

Document (JSON-like)

```
{  
  "id": 1,  
  "name": "Asha",  
  "mark": 85,  
  "courses": ["Math", "CS"]  
}
```

Great when data is naturally nested.
Eg: MongoDB, CouchDB

Key-Value

```
"student:1" → { ... }  
"student:2" → { ... }
```

Great for super-fast lookups by key.
Eg: Redis

Relational databases

Tables + IDs + relationships

Example: Users & Orders

Users

id	name
1	Asha
2	Rahul

Orders

id	user_id
501	1
502	1
503	2

Key ideas:

- “id” uniquely identifies a row (primary key)
- user_id links Orders → Users (relationship)

Why this is powerful


You can ask questions across tables.

Example question:
“Show order ids with user names”

Joins (concept)

Combine related rows from multiple tables

A join answers questions that span tables by matching IDs.



Takeaway: Relational databases shine when you need relationships + cross-table questions.

Non-relational (NoSQL) databases

Data stored in shapes beyond tables

NoSQL is a family of storage models, often chosen to match access patterns.

Document DB

Store a whole “thing” as a JSON-like document.

```
{  
  "id": 1,  
  "name": "Asha",  
  "orders": [501, 502]  
}
```

Good fit when your app typically fetches the whole object at once.

Key-Value store

A huge map: key → value

```
"session:abc" → { userId: 42 }  
"student:1"    → { ... }
```

Good fit for caching, sessions, and ultra-fast lookups.

NoSQL types (quick map)

Different tools for different access patterns

Document

JSON-like documents.
Good for user profiles, catalogs.

Key-Value

key -> value.
Good for caching & sessions.

Wide-column

Optimized for huge write throughput.
Good for logs/time-series.

Graph DB

Nodes + edges + traversals.
Good for recommendations / relationships.

Key idea: pick the model that matches your most common queries and reads/writes.

SQL vs NoSQL (high-level)

Not “which is better,” but which fits the problem

SQL (commonly relational)

- Data in tables + relationships
- Joins are a first-class idea
- Strong correctness tooling (constraints, transactions)
- Great for reporting & cross-entity questions

NoSQL (non-relational models)

- Data in documents / key-value / graph / ...
- Often model around app access patterns
- Flexible structure (varies by system)
- Great when you mostly fetch by key / whole document

Two myth-busters:

- NoSQL ≠ “no schema” (structure can still exist)
- SQL ≠ “can’t scale” (many SQL DBs scale well)

Constraints

Rules the database enforces to keep data correct

Why constraints exist:

App code can have bugs. Constraints are guardrails at the data layer.

Common constraints (examples)

- UNIQUE: email must be unique
- NOT NULL: name cannot be empty
- CHECK: marks ≥ 0
- FOREIGN KEY: order.user_id must exist in users

Bad data attempt

Constraint: UNIQUE(email)

```
INSERT user(email="asha@x")  
INSERT user(email="asha@x")
```

Result: second insert is rejected

When to use SQL (relational)

Choose SQL when relationships + correctness are central

SQL is a strong default when:

- Your data has clear structure (tables make sense).
- Relationships matter (users <-> orders <-> payments).
- You need cross-entity questions (joins, reports), where entities can be users, orders, etc
- Correctness is critical (constraints, transactions)

Examples: Payments, e-commerce orders, HR systems, analytics/reporting tables

When to use NoSQL

Choose a model that matches how your app reads/writes

NoSQL is a strong fit when:

- Data is naturally nested (JSON-like objects).
- Schema changes frequently or varies across records.
- You mostly fetch by key or by whole document.
- You want a specialized model (cache, graph traversal, time-series).

Examples: User profiles, product catalogs, caching, logs/events, social graphs

Decision guide: SQL vs NoSQL

A repeatable checklist you can apply to any scenario

Ask these 4 questions (in order):

1) Do relationships & joins matter?

Often SQL

2) Do you need multi-step correctness (transactions)?

Often SQL

3) Is the data naturally nested JSON?

Often Document DB

4) Is it mostly key lookups / caching / managing user sessions?

Often Key-Value

Quick exercise (30 sec): pick SQL or NoSQL for each and justify with the checklist.

- A) E-commerce orders
- B) Cache for recent searches
- C) User profile page