

Lecture 53

Database Integrations Using MongoDB

Listing, search, indexes, and data modeling

Today's objective and agenda

Objective

- Enable listing and search for students using MongoDB.
- Make index behavior visible using concrete examples.
- Model courses and enrollments using references.

Agenda

- Warm-up quiz.
- Live coding: listing and search.
- Activity: courses and enrollments.
- Indexes in practice.
- Embedded documents and references.

Warm-up

Warm-up quiz

- Answer eight multiple choice questions.
- Reply format: 1 / 2 / 3 / 4

Question 1

In a Spring Boot app, what are the two main inputs needed to connect Spring Data MongoDB to a database?

- 1) A Dockerfile and a MongoDB Compass connection
- 2) A MongoDB starter dependency and a MongoDB URI in application properties
- 3) A global exception handler and a seeded dataset
- 4) A controller annotation and a repository implementation class

Reply format: 1 / 2 / 3 / 4

Answer 1

In a Spring Boot app, what are the two main inputs needed to connect Spring Data MongoDB to a database?

- 1) A Dockerfile and a MongoDB Compass connection
- 2) A MongoDB starter dependency and a MongoDB URI in application properties
- 3) A global exception handler and a seeded dataset
- 4) A controller annotation and a repository implementation class

Correct answer: 2

Why: The starter enables Mongo support, and the URI tells Spring where the database is and which database name to use.

Reply format: 1 / 2 / 3 / 4

Question 2

What is the main job of Spring MVC in our application?

- 1) It stores the data in MongoDB collections.
- 2) It maps incoming HTTP requests to controller methods and produces HTTP responses.
- 3) It builds the database indexes at runtime for all collections.
- 4) It generates test data automatically during application startup.

Reply format: 1 / 2 / 3 / 4

Answer 2

What is the main job of Spring MVC in our application?

- 1) It stores the data in MongoDB collections.
- 2) It maps incoming HTTP requests to controller methods and produces HTTP responses.
- 3) It builds the database indexes at runtime for all collections.
- 4) It generates test data automatically during application startup.

Correct answer: 2

Why: Spring MVC routes requests to controllers and helps build the response, while repositories handle database persistence.

Reply format: 1 / 2 / 3 / 4

Question 3

What does the DispatcherServlet do in a Spring MVC application?

- 1) It inserts documents into MongoDB when the app starts.
- 2) It creates unique indexes automatically without configuration.
- 3) It replaces the need for service classes in the application.
- 4) It is the front controller that routes requests to the correct controller method.

Reply format: 1 / 2 / 3 / 4

Answer 3

What does the DispatcherServlet do in a Spring MVC application?

- 1) It inserts documents into MongoDB when the app starts.
- 2) It creates unique indexes automatically without configuration.
- 3) It replaces the need for service classes in the application.
- 4) It is the front controller that routes requests to the correct controller method.

Correct answer: 4

Why: DispatcherServlet is the front controller that receives requests first and forwards them to the right controller handler.

Reply format: 1 / 2 / 3 / 4

Question 4

Which repository method best represents pagination support when listing students?

- 1) Page<Student> findAll(Pageable pageable)
- 2) Optional<Student> findById(String id)
- 3) void deleteAll()
- 4) List<Student> findAll()

Reply format: 1 / 2 / 3 / 4

Answer 4

Which repository method best represents pagination support when listing students?

- 1) Page<Student> findAll(Pageable pageable)
- 2) Optional<Student> findById(String id)
- 3) void deleteAll()
- 4) List<Student> findAll()

Correct answer: 1

Why: Pageable carries page, size, and sort, and Page returns page content plus metadata such as total pages and total items.

Reply format: 1 / 2 / 3 / 4

Question 5

Which derived query method name follows Spring Data naming rules and filters by a boolean field?

- 1) findByActiveTrue()
- 2) searchStudentsWhereActivelstTrue()
- 3) getActiveStudents()
- 4) findStudentsActiveTrue()

Reply format: 1 / 2 / 3 / 4

Answer 5

Which derived query method name follows Spring Data naming rules and filters by a boolean field?

- 1) findByActiveTrue()
- 2) searchStudentsWhereActivelIsTrue()
- 3) getActiveStudents()
- 4) findStudentsActiveTrue()

Correct answer: 1

Why: Spring Data parses method names that follow the findBy + field + operator pattern.

Reply format: 1 / 2 / 3 / 4

Question 6

When fetching a student by id, what is the cleanest way for the repository layer to represent "not found"?

- 1) Return null and let the controller handle it.
- 2) Return Optional<Student> and let the service decide how to respond.
- 3) Return an empty string and let the service interpret it.
- 4) Return ResponseEntity<Student> from the repository.

Reply format: 1 / 2 / 3 / 4

Answer 6

When fetching a student by id, what is the cleanest way for the repository layer to represent "not found"?

- 1) Return null and let the controller handle it.
- 2) Return Optional<Student> and let the service decide how to respond.
- 3) Return an empty string and let the service interpret it.
- 4) Return ResponseEntity<Student> from the repository.

Correct answer: 2

Why: Optional cleanly represents presence or absence, and the service converts absence into a consistent API response such as 404.

Reply format: 1 / 2 / 3 / 4

Question 7

When two requests try to create the same email at nearly the same time, which layer provides the final guarantee that duplicates cannot be stored?

- 1) Controller validation annotations
- 2) A service-layer pre-check only
- 3) The database unique index
- 4) MongoDB Compass

Reply format: 1 / 2 / 3 / 4

Answer 7

When two requests try to create the same email at nearly the same time, which layer provides the final guarantee that duplicates cannot be stored?

- 1) Controller validation annotations
- 2) A service-layer pre-check only
- 3) The database unique index
- 4) MongoDB Compass

Correct answer: 3

Why: A unique index is enforced by MongoDB during the write, even when timing causes two requests to pass service checks.

Reply format: 1 / 2 / 3 / 4

Question 8

Which design best fits these two requirements?

Requirement A: a student has address fields.

Requirement B: an enrollment links one student to one course.

- 1) Store enrollments as separate documents that keep studentId and courseId.
- 2) Store everything in one collection and distinguish using a type field.
- 3) Store both address and enrollments inside the student document.
- 4) Store enrollments inside the course document.

Answer 8

Which design best fits these two requirements?

Requirement A: a student has address fields.

Requirement B: an enrollment links one student to one course.

- 1) Store enrollments as separate documents that keep studentId and courseId.
- 2) Store everything in one collection and distinguish using a type field.
- 3) Store both address and enrollments inside the student document.
- 4) Store enrollments inside the course document.

Correct answer: 1

Why: Enrollments are relationship records, so storing them as separate documents keeps the model clean and scalable.

Reply format: 1 / 2 / 3 / 4

Implement: Listing & Search

Live coding goal

- Enable a list endpoint that fetches students from MongoDB with filtering, sorting, and pagination.
- Reuse StudentPageResponse so the response returns items and page metadata together.
- Implement changes step by step so behavior is visible after each change.

Files that will be updated

- StudentController for the list endpoint and query parameters.
- StudentService for filtering logic and Pageable construction.
- StudentRepository for DB-backed queries.
- StudentPageResponse for the final response shape.

List endpoint behavior

- The endpoint accepts query parameters that control which students are returned and in what order.
- The endpoint returns a page response so the client can request the next page reliably.

Supported query parameters

- page and size control pagination.
- sortBy and sortDir control sorting.
- active filters by active status when provided.
- q searches by name or email when provided.

Example requests

```
/students?page=0&size=10&sortBy=name&sortDir=asc  
/students?active=true&q=ar&page=0&size=5&sortBy=name&sortDir=asc
```

Your turn!

Activity

Goal: Add the core domain building blocks for courses and enrollments.

- Create two new MongoDB document models: Course and Enrollment.
- Create two new repositories: CourseRepository and EnrollmentRepository.
- Keep the implementation minimal and consistent with the existing Student model and repository.
- Note: you are not required to do anything in MongoDB Compass for this activity.

Collections to be created

- courses
- enrollments

Requirement: Course document

- A course is stored as one document in the courses collection.
- Each course needs fields that support searching and display.

Required fields

- id as String, mapped to MongoDB _id.
- title as String.
- code as String, unique.
- active as boolean.

```
{  
  "_id": "65b9a0c2a9c7e24e0f4e20a1",  
  "title": "Spring Boot Foundations",  
  "code": "SB101",  
  "active": true  
}
```

Requirement: Enrollment document

- An enrollment is stored as one document in the enrollments collection.
- An enrollment connects one student to one course.
- Enrollment stores references so student and course remain independent documents.

Required fields

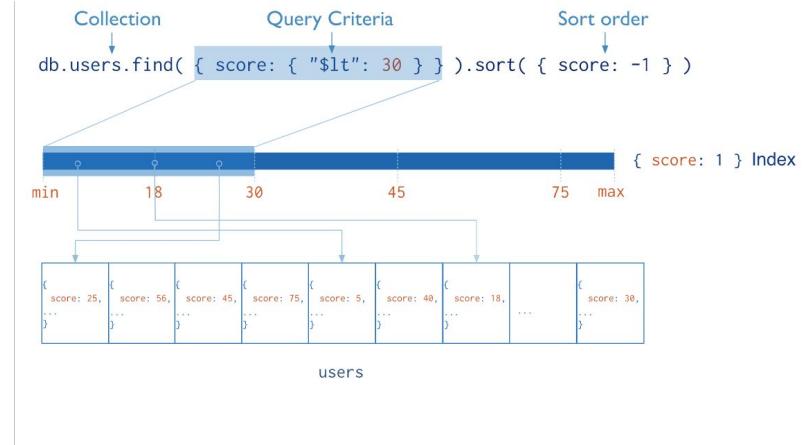
- id as String, mapped to MongoDB _id.
- studentId as String.
- courseId as String.
- enrolledAt as [Instant](#).
- status as String, for example ACTIVE or CANCELLED.

```
{  
  "_id": "65b9a13aa9c7e24e0f4e20b2",  
  "studentId": "65b7f3b2a9c7e24e0f4e1a91",  
  "courseId": "65b9a0c2a9c7e24e0f4e20a1",  
  "enrolledAt": "2026-02-01T10:15:00Z",  
  "status": "ACTIVE"  
}
```

Indexing

Recall: why do indexes matter in an API?

- An index exists to make reads faster by avoiding unnecessary scanning of documents.
- Without an index, MongoDB may need to examine many documents to find matches.
- With an index, MongoDB can jump directly to matching keys and fetch only relevant documents.
- Indexes can also enforce rules such as uniqueness when configured as unique.



Sample: Dataset for index demonstrations

- The examples assume the students collection stores fields like name, email, active, and city.
- The query examples (in upcoming slides) assume many students exist, not just a handful.

```
{ "_id": "65b7f3...a91", "name": "Rahul Kumar", "email": "rahul.k@gmail.com", "active": true, "city": "Bengaluru" }
{ "_id": "65b7f3...a92", "name": "Meera N", "email": "meera.n@gmail.com", "active": false, "city": "Chennai" }
{ "_id": "65b7f3...a93", "name": "Armaan Khan", "email": "armaan.k@gmail.com", "active": true, "city": "Bengaluru" }
{ "_id": "65b7f3...a94", "name": "Zoya Farhan", "email": "zoya.f@gmail.com", "active": true, "city": "Kochi" }
```

Single field index and the mental model

Query to support efficiently:

```
db.students.find({ city: "Bengaluru" })
```

Index that supports it:

```
db.students.createIndex({ city: 1 })
```

What MongoDB builds:

- MongoDB does not rearrange the documents inside the collection.
- MongoDB builds a separate index structure where keys are stored in sorted order.
- Each index key points to the documents that contain that value.

Index on city

```
"Bengaluru"  -> [ 65b7f3...a91, 65b7f3...a93 ]  
"Chennai"    -> [ 65b7f3...a92 ]  
"Kochi"       -> [ 65b7f3...a94 ]
```

Unique index as a database rule

- A unique index prevents two documents from having the same value for the indexed field.
- MongoDB enforces this rule during insert and update operations.

Example: email must be unique

```
db.students.createIndex({ email: 1 }, { unique: true })
```

Why this matters:

- Application code can try to prevent duplicates, but what if a query inserting duplicates is executed via MongoDB Compass?
- The database rule also protects data integrity when two requests are processed very close together.
Eg:

- Two requests can be processed concurrently (often on different server threads).
- If the app does check-then-insert (exists check, then save), both requests may pass the check before either inserts.
- Without a unique index, both inserts can succeed -> duplicates.
- With a unique index, one insert succeeds and the other fails with a [DuplicateKeyException](#) (which the app should catch and translate to a 409).

Duplicate email scenario described as a timeline

- Two create requests arrive close together for the same email.
- Request A checks and does not see the email yet.
- Before request A finishes saving, request B also checks and does not see the email yet.
- Both requests attempt to insert a new student.

What happens next:

- Without a unique index, both inserts can succeed and duplicates can be stored.
- With a unique index, MongoDB accepts the first insert and rejects the second insert.

Compound index and why order matters

- A compound index is built on more than one field.
- MongoDB stores the combined keys in sorted order.
- The first field creates the primary grouping, and the second field creates the secondary grouping.

Example compound index:

```
db.students.createIndex({ active: 1, city: 1 })
```

Index on (active, city)

```
(false, "Chennai")    -> [ 65b7f3...a92 ]  
(true,  "Bengaluru") -> [ 65b7f3...a91, 65b7f3...a93 ]  
(true,  "Kochi")     -> [ 65b7f3...a94 ]
```

- This index supports filtering by active efficiently.
- It also supports filtering by active and city together.
- It is not the best fit for queries that only filter by city.

Proving index usage beyond Compass

- Compass can show which indexes exist.
- `explain("executionStats")` shows which index MongoDB actually used for a query.

Command to run in mongosh:

```
db.students.find({ city: "Bengaluru" }).explain("executionStats")
```

What to look for:

- The plan indicates an index-based scan instead of a full collection scan.
- The execution stats show fewer examined documents when the index is used.

Embedded Docs & References

Recap: Modelling relationships in MongoDB

- MongoDB supports two common relationship patterns.
- You either embed related data inside one document, or store references between documents.

Embedding

- Store a smaller related object inside the parent document.
- Use this when the data is owned by the parent and is usually read together.

Referencing

- Store the related entity as its own document and store its id in another document.
- Use this when the related entity is shared, grows independently, or is queried separately.

Embedded example: Address inside Student

- Address fits well as embedded data because it belongs to a student.
- Address is usually read and updated together with the student.

A Student Document:

```
{  
  "_id": "65b7f3b2a9c7e24e0f4e1a91",  
  "name": "Rahul Kumar",  
  "email": "rahul.k@gmail.com",  
  "active": true,  
  "address": {  
    "line1": "123 MG Road",  
    "city": "Bengaluru",  
    "state": "Karnataka",  
    "pincode": "560001"  
  }  
}
```

Practical benefit:

- A single read returns the student and address together.
- No join-style lookup is required.

Reference example: Enrollment links Student and Course

- Enrollment should be a separate document because it is a relationship record.
- Many students can enroll into many courses.
- Enrollment data can grow independently and is often queried on its own.

```
{  
  "_id": "65b9a13aa9c7e24e0f4e20b2",  
  "studentId": "65b7f3b2a9c7e24e0f4e1a91",  
  "courseId": "65b9a0c2a9c7e24e0f4e20a1",  
  "enrolledAt": "2026-02-01T10:15:00Z",  
  "status": "ACTIVE"  
}
```

Practical benefit:

- Student and course documents remain small and stable.
- Enrollment queries do not force student documents to grow endlessly.

How this maps to our API and repositories

- Student remains the main API resource.
- Course and Enrollment enable relationship-based features.

Repository usage examples

- EnrollmentRepository.findByStudentId(studentId) lists what a student is enrolled in.
- EnrollmentRepository.findByCourseId(courseId) lists who is enrolled in a course.
- CourseRepository.findByCodeIgnoreCase(code) resolves courses by a human-readable key.

Next extension:

- Add endpoints that read enrollments and expand them into student and course details.

That's a wrap!